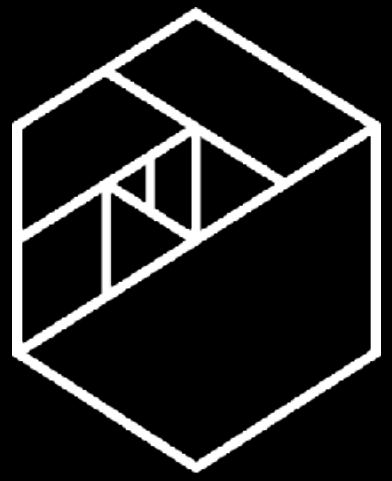

Review of Machine Learning

What is Machine Learning?

Machine learning allows computers to learn and infer from data.

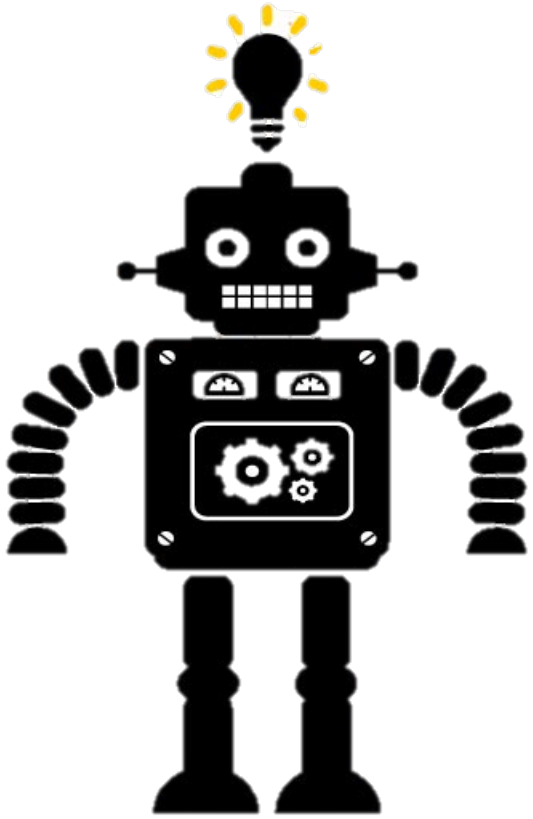

Types of Machine Learning

Supervised	data points have known outcome
Unsupervised	data points have unknown outcome

Types of Supervised Learning

Regression	data points have continuous outcome
Classification	data points have categorical outcome

Machine Learning Vocabulary

- **Target:** predicted category or value of the data (column to predict)
- **Features**: properties of the data used for prediction (non-target columns)
- **Example:** a single data point within the data (one row)
- **Label:** the target value for a single data point

Machine Learning Vocabulary (Synonyms)

- **Target:** Response, Output, Dependent Variable, Labels
- **Features**: Predictors, Input, Independent Variables, Attributes
- **Example:** Observation, Record, Instance, Datapoint, Row
- **Label:** Answer, y-value, Category

Supervised Learning Overview

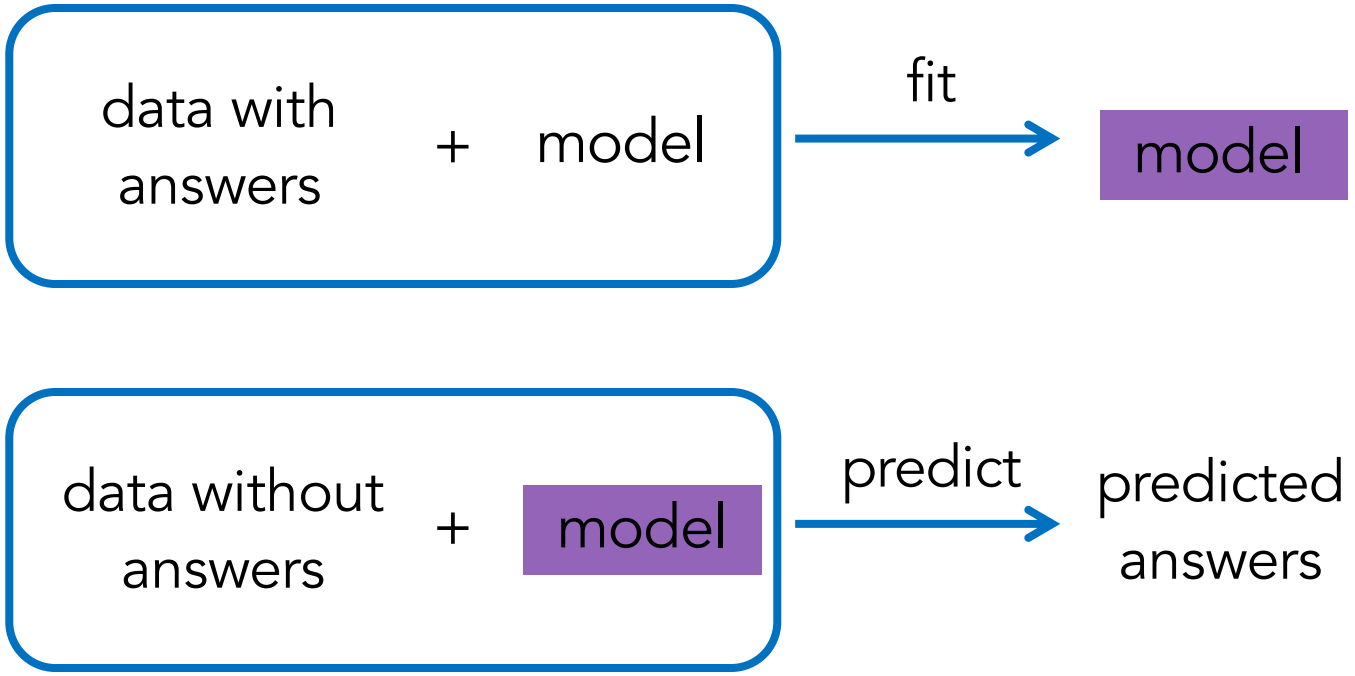

Regression: Numeric Answers

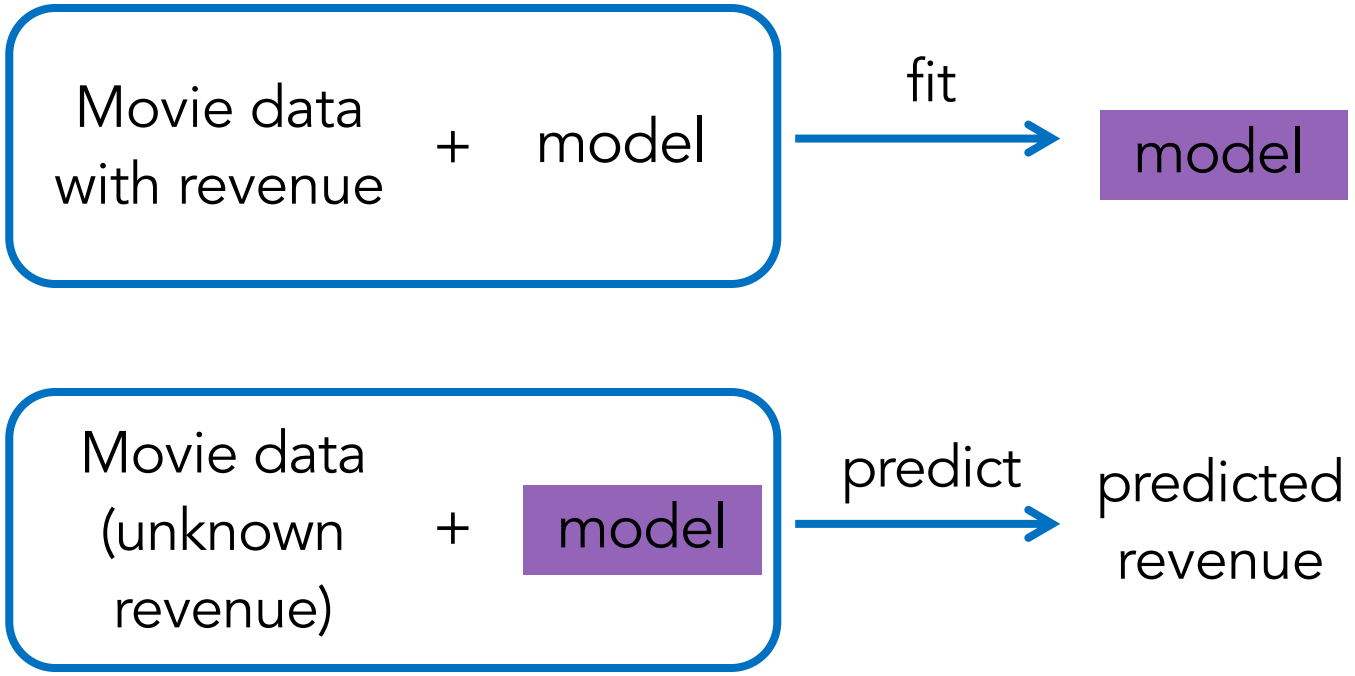

Classification: Categorical Answers

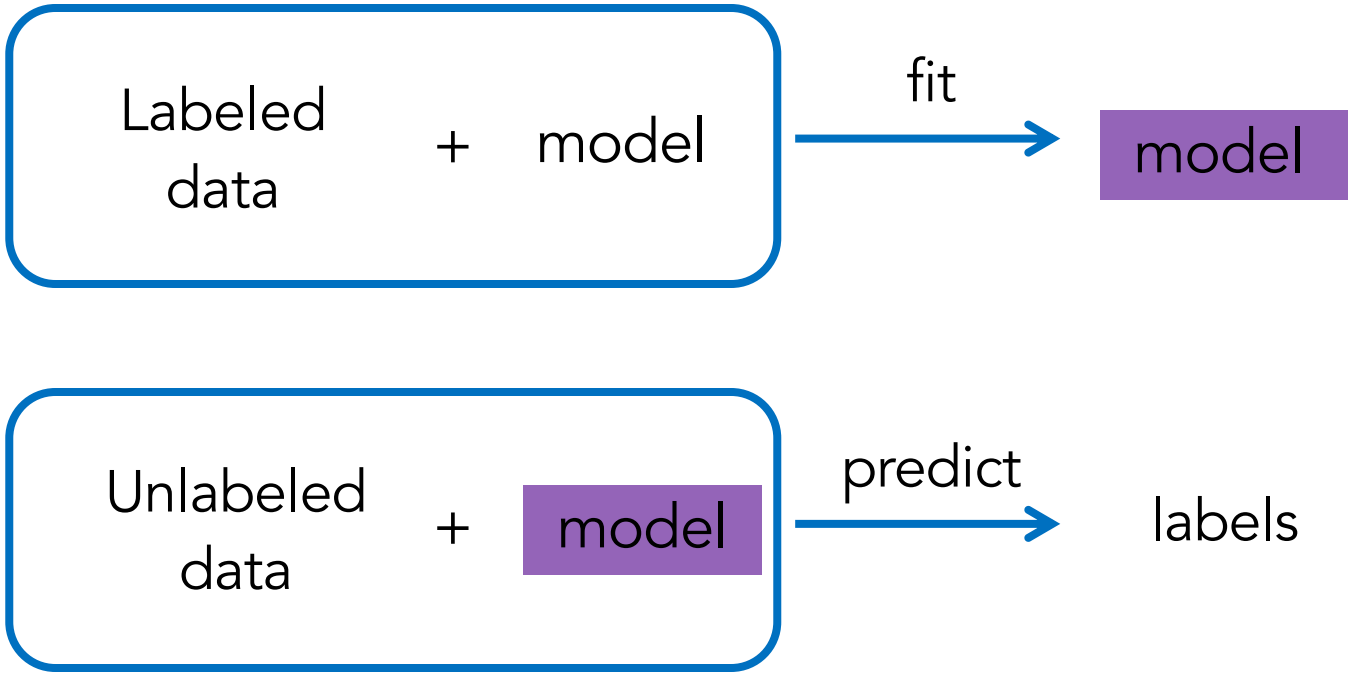

Classification: Categorical Answers

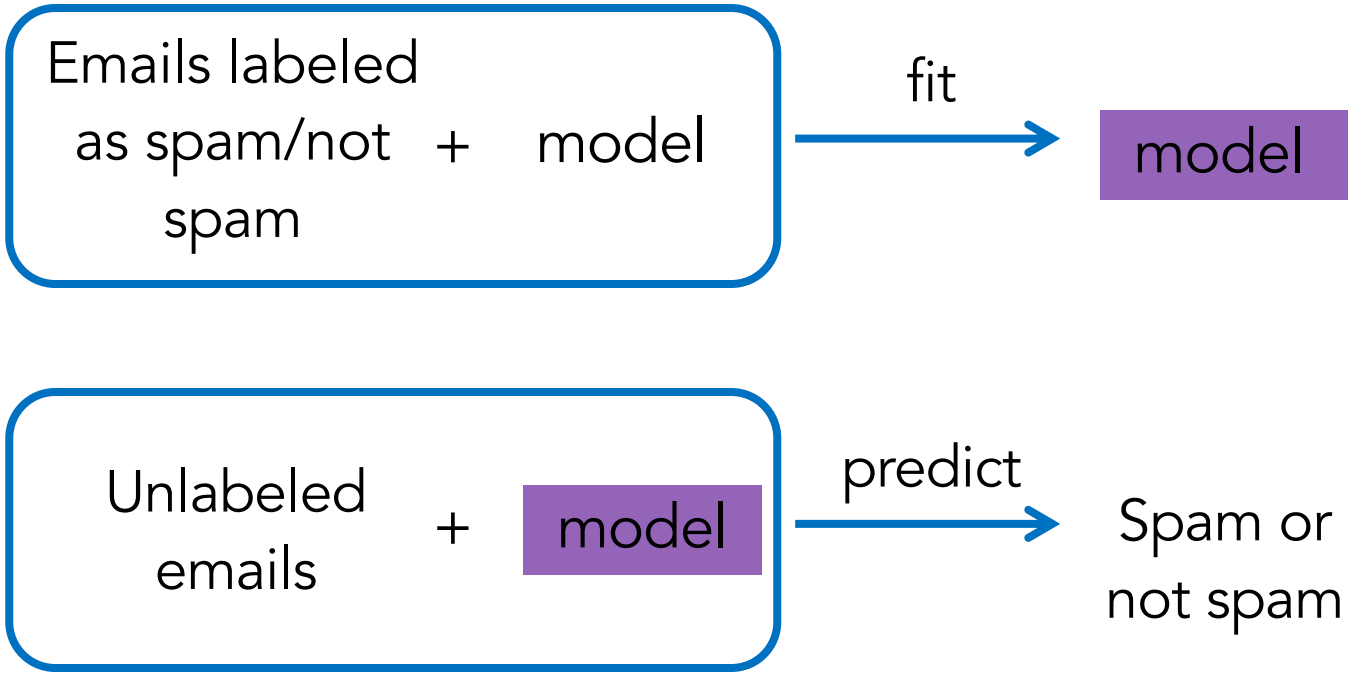

Two Types of Classification Predictions

- **Hard Prediction:** Predict a single category for each instance.
- **Probability Prediction:** Assign a probability distribution across the classes to each instance.

Metrics for Classification

- **Hard Prediction:** Accuracy, Precision, Recall (Sensitivity), Specificity, F1 Score
- **Probability Prediction:** Log-loss (aka Cross-Entropy), Brier Score, AUC (ROC), Precision-Recall Curves

Metrics for Regression

- Root Mean Square Error (RMSE)

$$RMSE = \sqrt{\frac{1}{n}\sum_{i=1}^{n}(y_i - \hat{y}_i)^2}$$

- Mean Absolute Deviation

$$MAD = \frac{1}{n}\sum_{i=1}^{n}|y_i - \hat{y}_i|$$

Fitting Training and Test Data

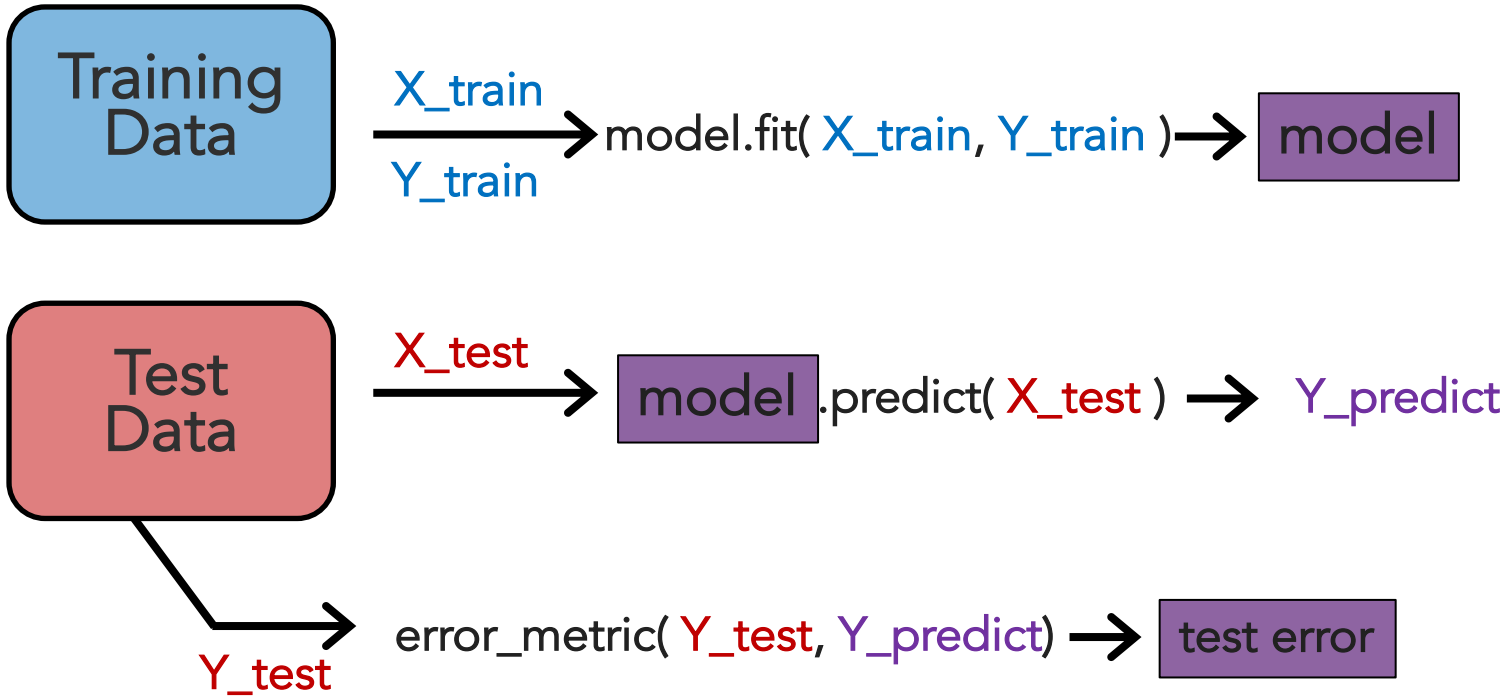

Using Training and Test Data

Training Data

fit the model

Test Data

measure performance
- predict label with model
- compare with actual value
- measure error

Using Training and Test Data

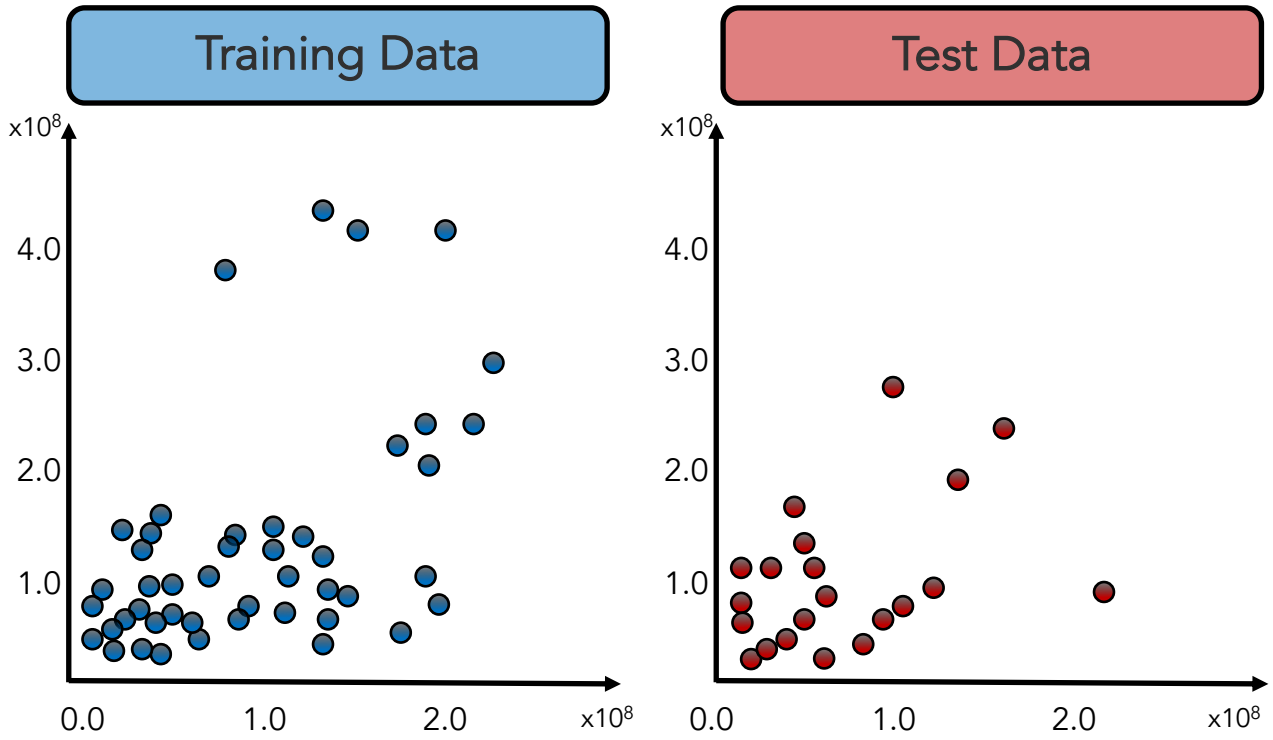

Using Training and Test Data

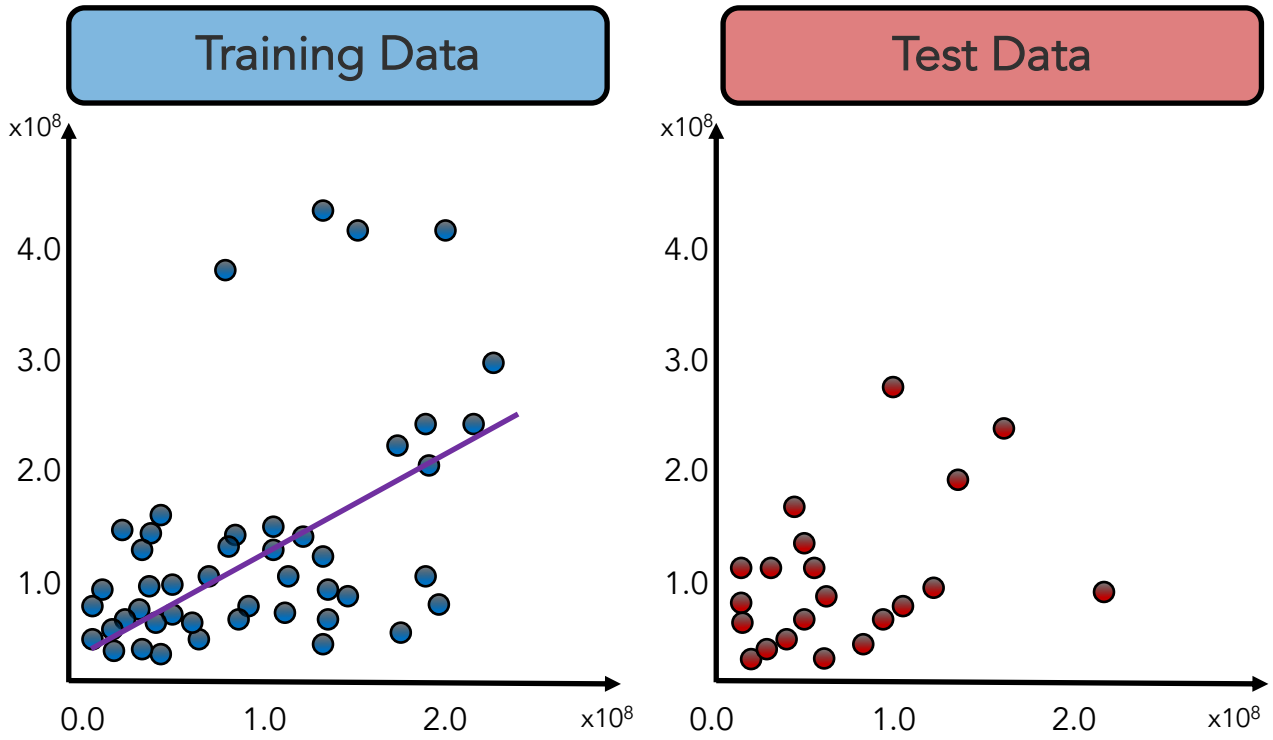

Using Training and Test Data

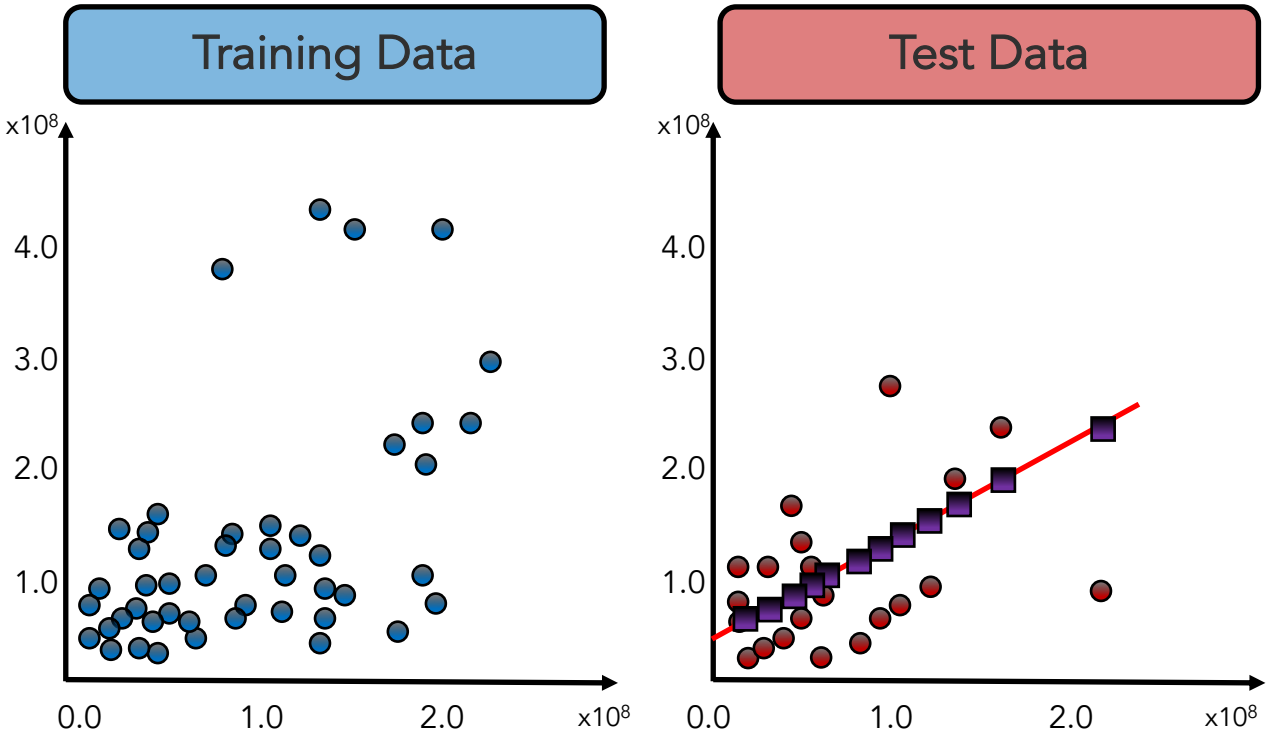

Make predictions

Using Training and Test Data

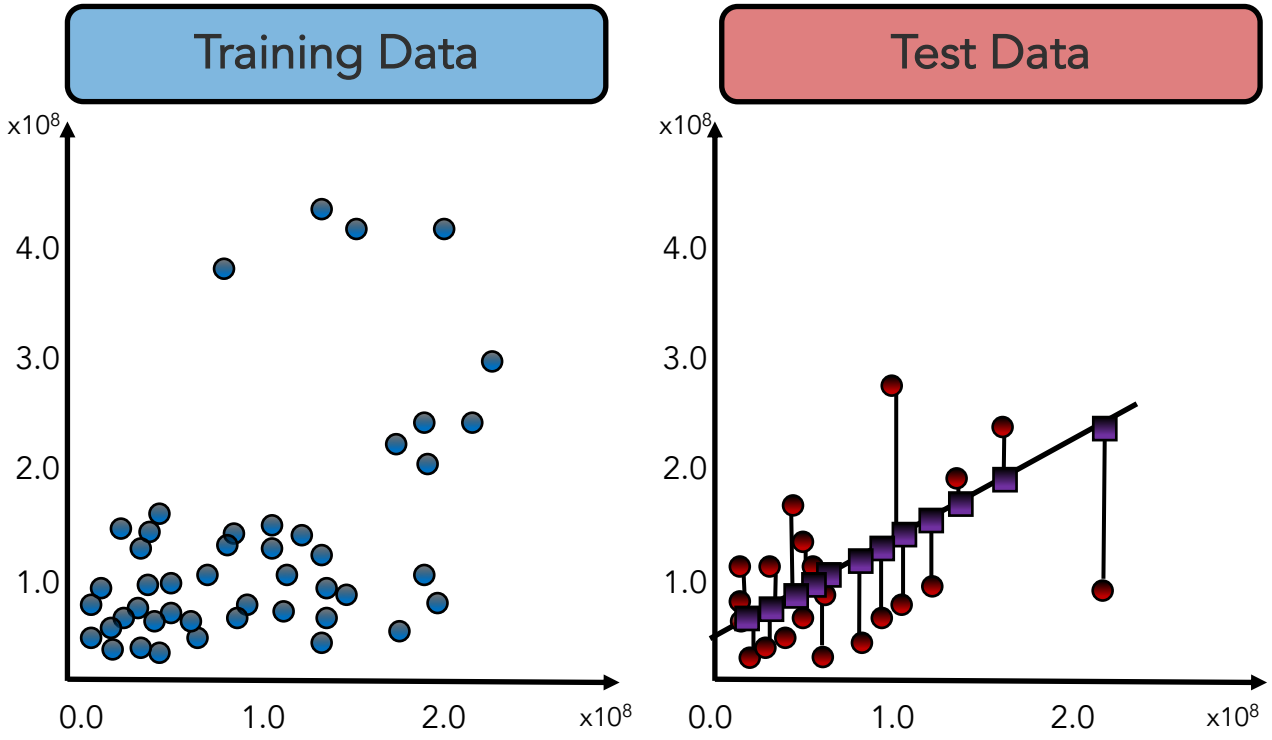

Measure error

How Well Does the Model Generalize?

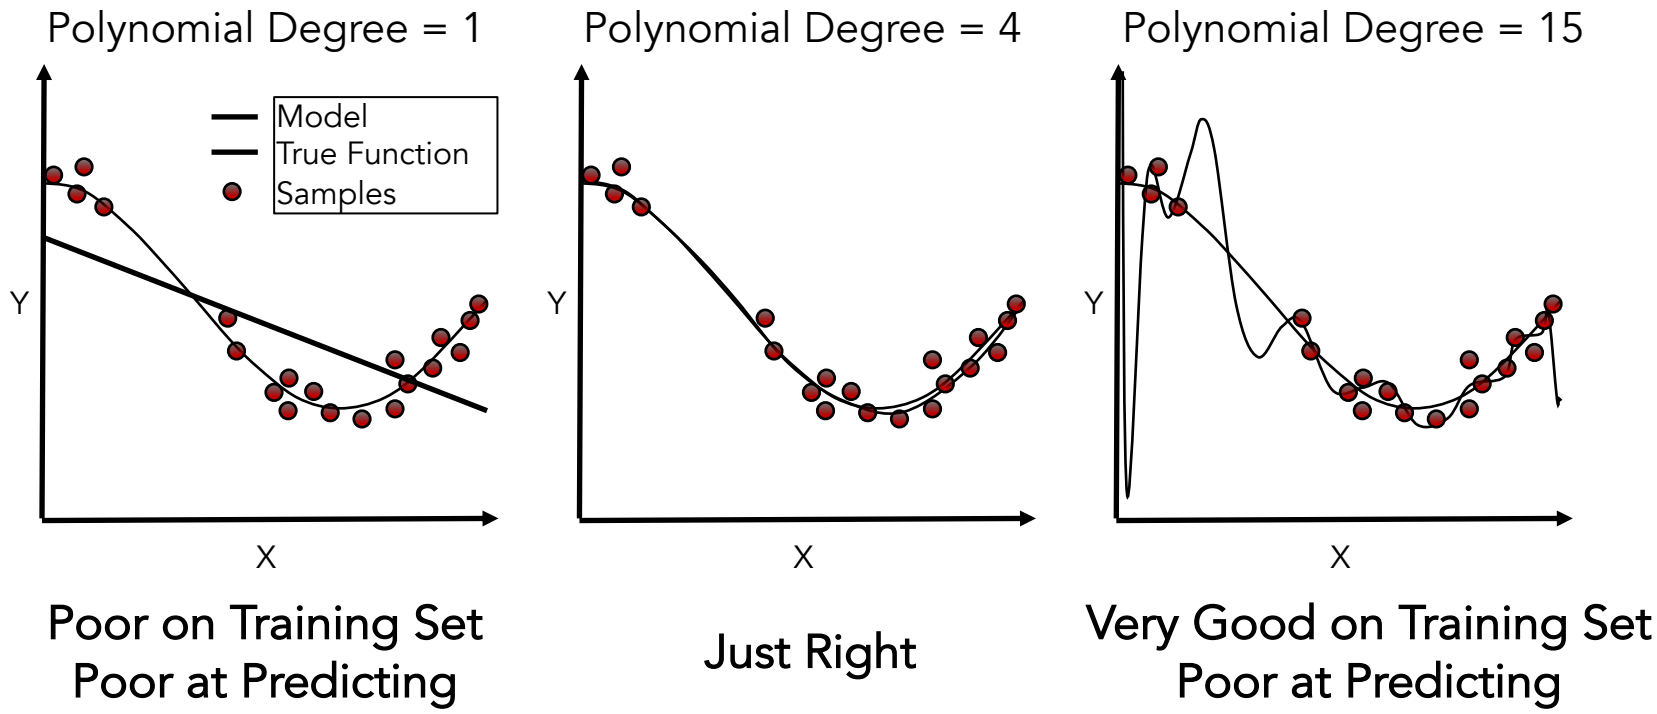

Poor on Training Set
Poor at Predicting

Just Right

Very Good on Training Set
Poor at Predicting

Underfitting vs Overfitting

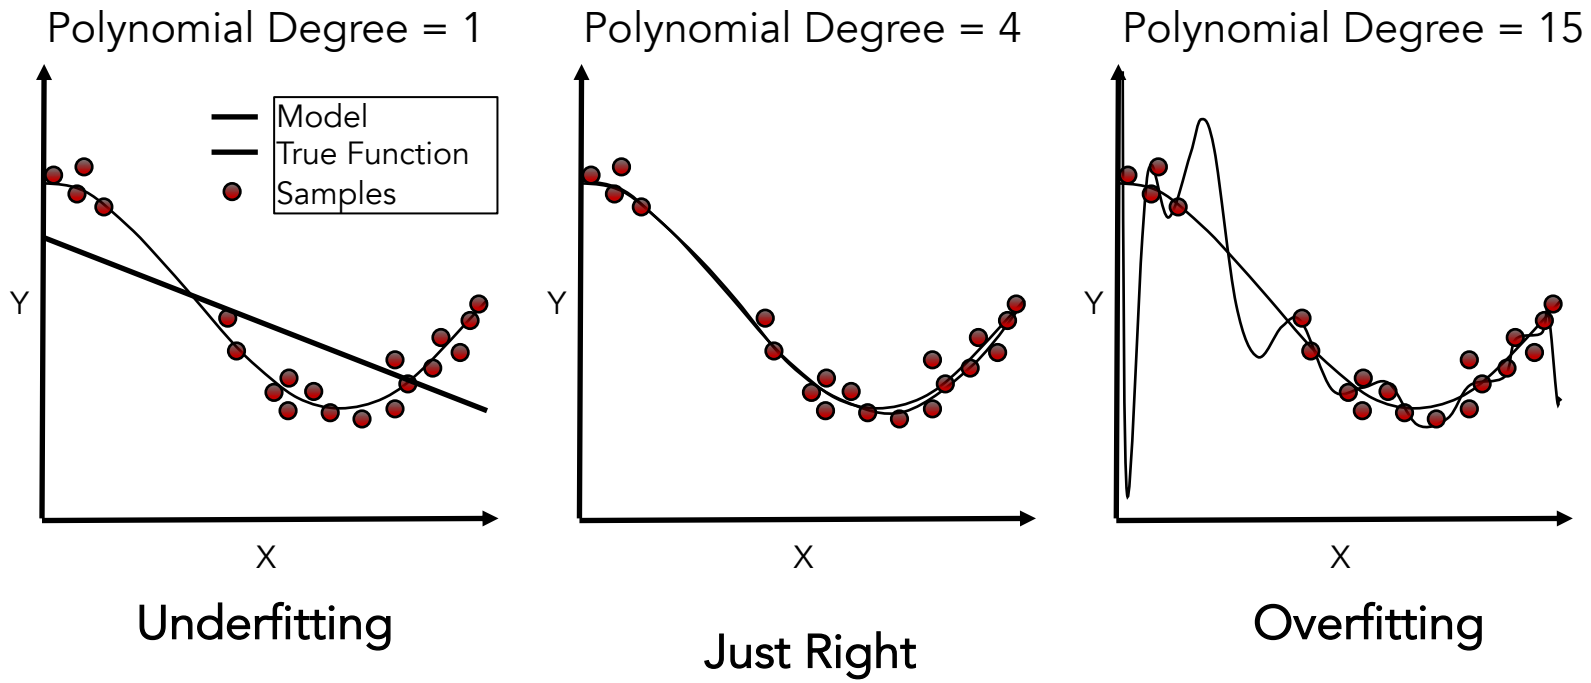

Bias/Variance Tradeoff

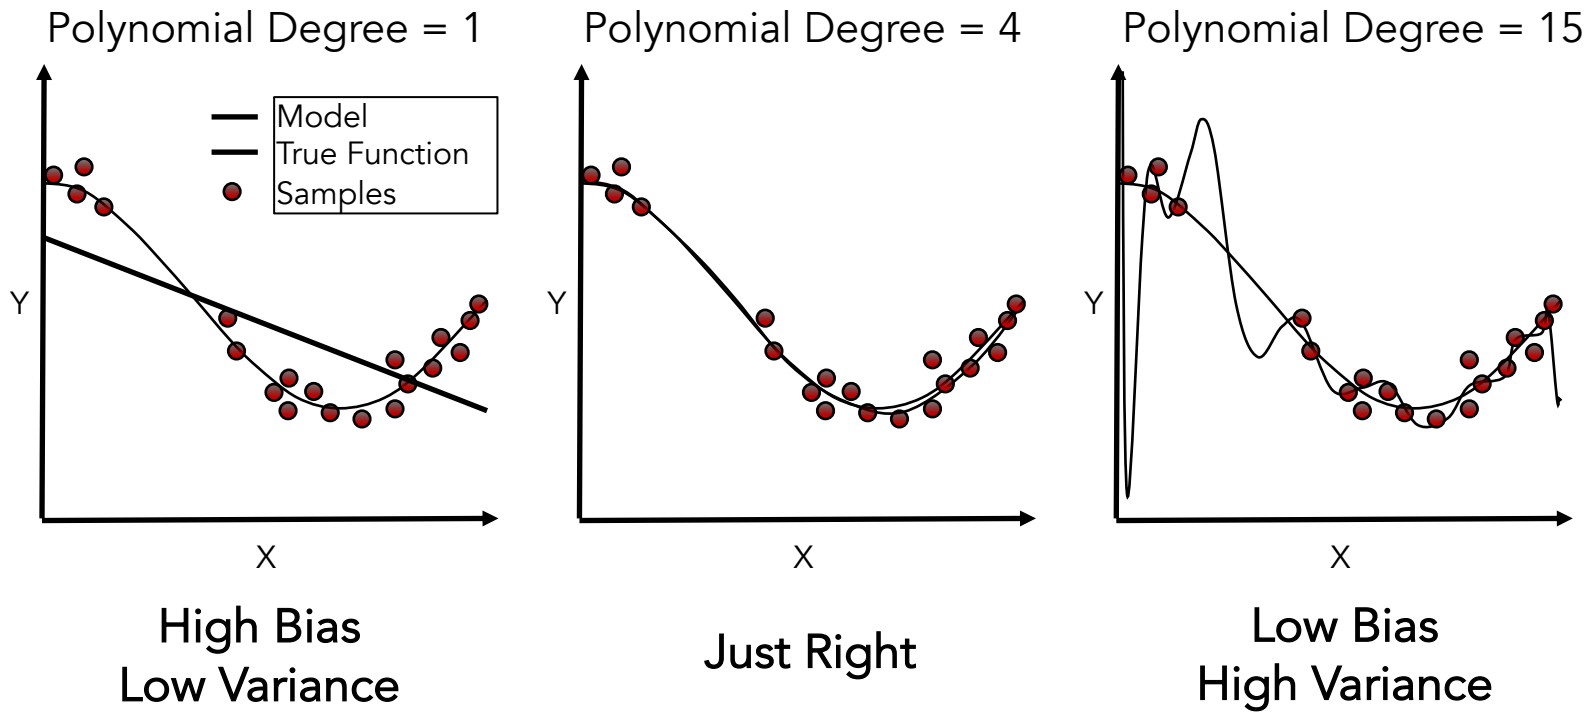

Questions?

METIS